

2026 RESEARCH PREVIEW

Agent Braid

Open tooling outlook

Portable, evidence-carrying concurrency analysis
for multi-agent systems

Internal evidence · Independent validation pending



The open-tooling thesis

What the ecosystem already does

Orchestration decides who acts

Protocols connect agents and tools

Observability records what happened

The layer under test

Which operations may overlap?

Which operations must be ordered?

What evidence supports the decision?

Agent Braid tests a conservative, vendor-neutral evidence layer before any execution runtime exists.

Observed scale signals, not capturable demand

180M+

developers reported by GitHub
in Octoverse 2025

4.3M

AI projects reported by GitHub;
projects are not organizations


1M+

OpenAI business customers;
one vendor signal, not global TAM

These units differ and are never summed. They justify investigation only; none measures organizations willing to adopt or pay for Agent Braid.

Validation priority: organization counts · conflict incidence · adoption willingness · budget ownership

Evidence between orchestration and execution




Evidence-gated path to international impact

Software first. Platforms after stable contracts. Enterprise execution only behind defensible controls.



M3 braid semantics remains an independent research track. A negative braid result cannot invalidate useful, sound concurrency engineering.

Opportunity proxies and community targets



18-month target	Low	Base	High
Integrations	2	4	7
Independent reproductions	3	10	25
Repeat contributors	5	15	40

Scenario inputs are assumptions. Validate conflict incidence, adoption willingness, integration cost and budget ownership before commercial interpretation.

TAM/SAM are organization proxies; SOM is community adoption. The model never adds incompatible units.